

Estimating quantum speedups for lattice sieves

Martin R. Albrecht¹, Vlad Gheorghiu², *Eamonn W. Postlethwaite*¹, John M. Schanck²

¹Information Security Group, Royal Holloway, University of London,

²Institute for Quantum Computing, University of Waterloo, Canada

What are we doing? We are

- trying to understand 'quantum lattice sieves' non-asymptotically,

What are we doing? We are

- trying to understand ‘quantum lattice sieves’ non-asymptotically,
- trying to glance behind the query model – e.g. no longer counting Grover oracle queries,

What are we doing? We are

- trying to understand ‘quantum lattice sieves’ non-asymptotically,
- trying to glance behind the query model – e.g. no longer counting Grover oracle queries,
- trying to understand the quantum overhead of these sieves, and compare to their classical variants.

How to do this? We need to

- decide exactly what a query to our oracle is,



How to do this? We need to

- decide exactly what a query to our oracle is,
- build a circuit for it,



How to do this? We need to

- decide exactly what a query to our oracle is,
- build a circuit for it,
- describe some cost metrics to optimise controllable parameters under,



How to do this? We need to

- decide exactly what a query to our oracle is,
- build a circuit for it,
- describe some cost metrics to optimise controllable parameters under,
- build some software to perform this optimisation.



Why is this interesting? (Good question) because

- a great deal of cryptography, some close to standardisation, uses lattice based assumptions,


Why is this interesting? (Good question) because

- a great deal of cryptography, some close to standardisation, uses lattice based assumptions,
- classically it is lattice sieves that currently power the best cryptanalysis,

Why is this interesting? (Good question) because


- a great deal of cryptography, some close to standardisation, uses lattice based assumptions,
- classically it is lattice sieves that currently power the best cryptanalysis,
- what if a large fault tolerant quantum computer appeared at CWI tomorrow?

What: lattices




$$\Lambda = \text{Span}_{\mathbb{Z}}(b_0, \dots, b_{d-1}), B = \{b_0, \dots, b_{d-1}\} \subset \mathbb{R}^d \text{ a basis}$$

What: lattices




SVP: find $v \in \Lambda \setminus \{0\}$ such that $\|v\|_2 \leq \|w\|_2$ for all $w \in \Lambda \setminus \{0\}$

What: lattices




α -HSVP: find $v \in \Lambda \setminus \{0\}$ such that $\|v\|_2 \leq \alpha \cdot \text{vol}(\Lambda)^{1/d}$

What: lattice sieves



We can cheaply sample a long lattice vector with uniform direction in some thin annulus.


What: lattice sieves



We can cheaply sample a long lattice vector with uniform direction in some thin annulus.

Heuristic sieves are analysed as if we have many uniform vectors in S^{d-1} .

What: lattice sieves




We can cheaply sample a long lattice vector with uniform direction in some thin annulus.

Heuristic sieves are analysed as if we have many uniform vectors in S^{d-1} .


In this model $u - v$ is shorter (within the circle) iff $\theta(u, v) < \pi/3$.

What: lattice sieves (high dimensions)



As the dimension grows the distribution of $\theta(u, v)$ becomes concentrated around $\pi/2$.


What: lattice sieves (high dimensions)



As the dimension grows the distribution of $\theta(u, v)$ becomes concentrated around $\pi/2$.

To find sufficiently many pairs with $\theta(u, v) < \pi/3$ we require $\exp(d)$ vectors; most reductions will have $\theta(u, v) \approx \pi/3$.

What: lattice sieves (high dimensions)




As the dimension grows the distribution of $\theta(u, v)$ becomes concentrated around $\pi/2$.

To find sufficiently many pairs with $\theta(u, v) < \pi/3$ we require $\exp(d)$ vectors; most reductions will have $\theta(u, v) \approx \pi/3$.

We begin anew with some thin annulus of vectors an $\varepsilon \in (0, 1)$ factor shorter.

What: lattice sieves (bucketing)



Calculating $\theta(u, v)$ is effectively an inner product, the number of which we want to minimise.

Lattice sieves therefore bucket vectors in various manners and check $\theta(u, v)$ only within these buckets.

One can also filter further within buckets (spoiler: we do this).

What: different lattice sieves

Sieve (NNS subroutine) ¹	$\log_2 \text{time}_C$	$\log_2 \text{time}_Q$
NV style [NV08]	$0.415d$	$0.311d$
RandomBucket [BGJ15, ADH ⁺ 19]	$0.349d$	$0.301d$
ListDecoding [BDGL16, DSvW21]	$0.292d$	$0.265d$

¹All complexities are missing $+o(d)$ terms.

What: different lattice sieves

Sieve (NNS subroutine) ¹	$\log_2 \text{time}_C$	$\log_2 \text{time}_Q$
NV style [NV08]	$0.415d$	$0.311d$
RandomBucket [BGJ15, ADH ⁺ 19]	$0.349d$	$0.301d$
ListDecoding [BDGL16, DSvW21]	$0.292d$	$0.265d$

The quantum variants of these sieves use Grover's search algorithm to instantiate the search for reducing pairs (within buckets, when appropriate).

All require exponential space, $2^{\Theta(d)}$.

¹All complexities are missing $+o(d)$ terms.

How: classical and quantum search

Let $[N] = \{1, \dots, N\}$ and $f: [N] \rightarrow \{0, 1\}$ be an unstructured predicate, with *roots*

$$\text{Ker}(f) = \{x: f(x) = 0\}.$$

How: classical and quantum search

Let $[N] = \{1, \dots, N\}$ and $f: [N] \rightarrow \{0, 1\}$ be an unstructured predicate, with *roots*

$$\text{Ker}(f) = \{x: f(x) = 0\}.$$

We can find a root

- classically by evaluating $f(1), \dots, f(N)$,
- quantumly by measuring $\mathbf{G}(f)^j \mathbf{D}|0\rangle$.

How: classical and quantum search

Let $[N] = \{1, \dots, N\}$ and $f: [N] \rightarrow \{0, 1\}$ be an unstructured predicate, with *roots*

$$\text{Ker}(f) = \{x: f(x) = 0\}.$$

We can find a root

- classically by evaluating $f(1), \dots, f(N)$,
- quantumly by measuring $\mathbf{G}(f)^j \mathbf{D}|0\rangle$.

If $|\text{Ker}(f)| \in o(N)$ then, to succeed with constant probability, we expect $O(N)$ queries to f classically, and $j \in O(\sqrt{N})$ queries to $\mathbf{G}(f)$ quantumly.

How: filtered search

Classically, a potentially cheaper way is to use a filter, some predicate

$$g: [N] \rightarrow \{0, 1\}, |\text{Ker}(g) \cap \text{Ker}(f)| \geq 1.$$

Then we can

How: filtered search

Classically, a potentially cheaper way is to use a filter, some predicate

$$g: [N] \rightarrow \{0, 1\}, |\text{Ker}(g) \cap \text{Ker}(f)| \geq 1.$$

Then we can

evaluate $g(1)$, if $g(1) = 0$ evaluate $f(1)$,

How: filtered search

Classically, a potentially cheaper way is to use a filter, some predicate

$$g: [N] \rightarrow \{0, 1\}, |\text{Ker}(g) \cap \text{Ker}(f)| \geq 1.$$

Then we can

evaluate $g(1)$, if $g(1) = 0$ evaluate $f(1)$, \dots , evaluate $g(N)$, if $g(N) = 0$ evaluate $f(N)$.

How: filtered search

Classically, a potentially cheaper way is to use a filter, some predicate

$$g: [N] \rightarrow \{0, 1\}, |\text{Ker}(g) \cap \text{Ker}(f)| \geq 1.$$

Then we can

evaluate $g(1)$, if $g(1) = 0$ evaluate $f(1)$, \dots , evaluate $g(N)$, if $g(N) = 0$ evaluate $f(N)$.


What makes a good filter? Cheaper than f to evaluate, and

$$\rho_f(g) = 1 - \frac{|\text{Ker}(f) \cap \text{Ker}(g)|}{|\text{Ker}(g)|}, \quad \eta_f(g) = 1 - \frac{|\text{Ker}(f) \cap \text{Ker}(g)|}{|\text{Ker}(f)|}$$

the false positive and negative rate, are both small.


How: filtered quantum search

Branching based on g is not possible within Grover's algorithm. However we can use *amplitude amplification* to achieve something conceptually similar. First, Grover's algorithm over g :




How: filtered quantum search

Branching based on g is not possible within Grover's algorithm. However we can use *amplitude amplification* to achieve something conceptually similar. First, Grover's algorithm over g :




How: filtered quantum search

Branching based on g is not possible within Grover's algorithm. However we can use *amplitude amplification* to achieve something conceptually similar. First, Grover's algorithm over g :




How: filtered quantum search

Branching based on g is not possible within Grover's algorithm. However we can use *amplitude amplification* to achieve something conceptually similar. First, Grover's algorithm over g :




How: filtered quantum search (amplitude amplification)

Amplitude amplification can replace \mathbf{D} with $\mathbf{A}_j = \mathbf{G}(g)^j \mathbf{D}$. Then amplitude amplification for the predicate $f \cap g$:




How: filtered quantum search (amplitude amplification)

Amplitude amplification can replace \mathbf{D} with $\mathbf{A}_j = \mathbf{G}(g)^j \mathbf{D}$. Then amplitude amplification for the predicate $f \cap g$:



How: filtered quantum search (amplitude amplification)

Amplitude amplification can replace \mathbf{D} with $\mathbf{A}_j = \mathbf{G}(g)^j \mathbf{D}$. Then amplitude amplification for the predicate $f \cap g$:



How: filtered quantum search (with partial information)

We give some technical results that (roughly) say, let

How: filtered quantum search (with partial information)

We give some technical results that (roughly) say, let

- g be a filter for predicate $f: [N] \rightarrow \{0, 1\}$,
- $P, Q, \gamma \in \mathbb{R}$ such that
 - $P/\gamma \leq |\text{Ker}(g)| \leq \gamma P$, and
 - $1 \leq Q \leq |\text{Ker}(f) \cap \text{Ker}(g)|$.

How: filtered quantum search (with partial information)

We give some technical results that (roughly) say, let

- g be a filter for predicate $f: [N] \rightarrow \{0, 1\}$,
- $P, Q, \gamma \in \mathbb{R}$ such that
 - $P/\gamma \leq |\text{Ker}(g)| \leq \gamma P$, and
 - $1 \leq Q \leq |\text{Ker}(f) \cap \text{Ker}(g)|$.

Then we can find a root of f with constant probability and a cost dominated by $\frac{\gamma}{2} \sqrt{N/Q}$ calls to $\mathbf{G}(g)$.

How: filtered quantum search (with partial information)

We give some technical results that (roughly) say, let

- g be a filter for predicate $f: [N] \rightarrow \{0, 1\}$,
- $P, Q, \gamma \in \mathbb{R}$ such that
 - $P/\gamma \leq |\text{Ker}(g)| \leq \gamma P$, and
 - $1 \leq Q \leq |\text{Ker}(f) \cap \text{Ker}(g)|$.

Then we can find a root of f with constant probability and a cost dominated by $\frac{\gamma}{2} \sqrt{N/Q}$ calls to $\mathbf{G}(g)$.

The idea: the cost of a Grover query encoding the filter, $\mathbf{G}(g)$, is the crucial quantity.

\Rightarrow specify g , design $\mathbf{G}(g)$, and understand (P, Q, γ) .

How: popcount is our filter

For lattice vectors u, v_1, \dots, v_N , the reduction predicate of u is


$$f_u: \{v_1, \dots, v_N\} \rightarrow \{0, 1\}, \quad f_u(v_i) = 0 \iff \langle u, v_i \rangle > \cos(\pi/3).$$

How: popcount is our filter

For lattice vectors u, v_1, \dots, v_N , the reduction predicate of u is

$$f_u: \{v_1, \dots, v_N\} \rightarrow \{0, 1\}, \quad f_u(v_i) = 0 \iff \langle u, v_i \rangle > \cos(\pi/3).$$

For the filter g we use ‘XOR and popcount’ [FBB⁺14], i.e. $g_u(\cdot) = \text{popcount}_{k,n}(u, \cdot)$.



$$\text{popcount}_{1,2}(u, v_0) = 1$$

$$\text{popcount}_{1,2}(u, v_1) = 0$$

$$\text{popcount}_{1,2}(u, v_2) = 0$$


$$(k, n) = (1, 2)$$

How: popcount is our filter

For lattice vectors u, v_1, \dots, v_N , the reduction predicate of u is

$$f_u: \{v_1, \dots, v_N\} \rightarrow \{0, 1\}, \quad f_u(v_i) = 0 \iff \langle u, v_i \rangle > \cos(\pi/3).$$

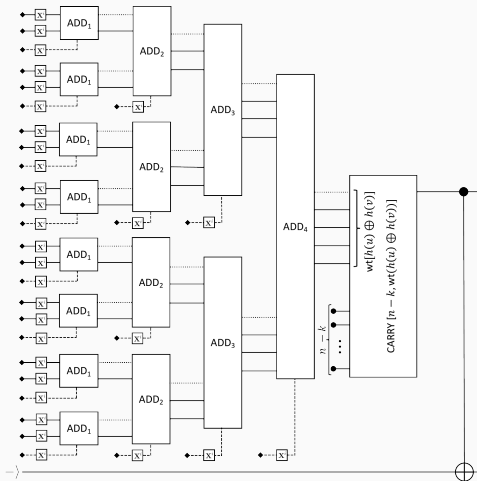
For the filter g we use ‘XOR and popcount’ [FBB⁺14], i.e. $g_u(\cdot) = \text{popcount}_{k,n}(u, \cdot)$.



$$h_i: \mathbb{R}^d \rightarrow \{0, 1\}$$

$$\text{check, over } \mathbb{Z}: \sum_i h_i(u) \oplus h_i(v_j) \leq k$$

How: circuits for $G(\text{popcount}_{k,n})$



Basically a (reversible) tree of in place quantum adders ending with a comparison.

How: a probabilistic study of popcount

Given i.i.d. uniform $\{h_i\}_{i=1}^n$, some threshold k , and pair (u, v) on S^{d-1} , let $P_{k,n}(u, v)$ be the probability the pair pass $\text{popcount}_{k,n}$. Then

How: a probabilistic study of popcount

Given i.i.d. uniform $\{h_i\}_{i=1}^n$, some threshold k , and pair (u, v) on S^{d-1} , let $P_{k,n}(u, v)$ be the probability the pair pass $\text{popcount}_{k,n}$. Then

$$\Pr[P_{k,n}(u, v)] = \sum_{i=0}^k \binom{n}{i} \cdot \left(\frac{\theta(u, v)}{\pi} \right)^i \cdot \left(1 - \frac{\theta(u, v)}{\pi} \right)^{n-i}.$$

How: a probabilistic study of popcount

Given i.i.d. uniform $\{h_i\}_{i=1}^n$, some threshold k , and pair (u, v) on S^{d-1} , let $P_{k,n}(u, v)$ be the probability the pair pass $\text{popcount}_{k,n}$. Then

$$\Pr[P_{k,n}(u, v)] = \sum_{i=0}^k \binom{n}{i} \cdot \left(\frac{\theta(u, v)}{\pi} \right)^i \cdot \left(1 - \frac{\theta(u, v)}{\pi} \right)^{n-i}.$$

Ultimately it is $\theta = \theta(u, v)$ that matters, so we consider $P_{k,n}(\theta)$.

How: a simple example

The pdf of two uniform $u, v \in S^{d-1}$ having $\theta(u, v) = \theta$ is

$$A_d(\theta) = C(d) \cdot \sin^{d-2}(\theta),$$

and the probability of u, v passing $\text{popcount}_{k,n}$ is then given by

$$\int_0^\pi P_{k,n}(\theta) \cdot A_d(\theta) \, d\theta.$$

How: a simple example

The pdf of two uniform $u, v \in S^{d-1}$ having $\theta(u, v) = \theta$ is

$$A_d(\theta) = C(d) \cdot \sin^{d-2}(\theta),$$

and the probability of u, v passing $\text{popcount}_{k,n}$ is then given by

$$\int_0^\pi P_{k,n}(\theta) \cdot A_d(\theta) \, d\theta.$$

For the false negative rate and the different bucketing strategies we integrate $P_{k,n}(\theta)$ over the relevant spherical sections.

How: cost metrics

Following [JS19] we measure the cost of running a quantum circuit in terms of the classical control required to run it. Here (G, D, W) are (gate count, depth, width) of a quantum circuit.

How: cost metrics

Following [JS19] we measure the cost of running a quantum circuit in terms of the classical control required to run it. Here (G, D, W) are (gate count, depth, width) of a quantum circuit.

- gates: quantum gates cost $\Theta(1) \xRightarrow{\text{total}} \Theta(G)$,

How: cost metrics

Following [JS19] we measure the cost of running a quantum circuit in terms of the classical control required to run it. Here (G, D, W) are (gate count, depth, width) of a quantum circuit.

- gates: quantum gates cost $\Theta(1) \xRightarrow{\text{total}} \Theta(G)$,
- depth-width: {quantum gates, identity wires} cost $\Theta(1) \xRightarrow{\text{total}} \Theta(DW)$,
- error: {quantum gates, identity wires} cost $\Theta(\log^2(DW)) \xRightarrow{\text{total}} \Omega(DW \log^2(DW))$.

How: cost metrics

Following [JS19] we measure the cost of running a quantum circuit in terms of the classical control required to run it. Here (G, D, W) are (gate count, depth, width) of a quantum circuit.

- gates: quantum gates cost $\Theta(1) \xrightarrow{\text{total}} \Theta(G)$,
- depth-width: {quantum gates, identity wires} cost $\Theta(1) \xrightarrow{\text{total}} \Theta(DW)$,
- error: {quantum gates, identity wires} cost $\Theta(\log^2(DW)) \xrightarrow{\text{total}} \Omega(DW \log^2(DW))$.


In particular we use the error correction model of Gidney–Ekerå [GE19] and the Clifford+ T gate set. We compliment it with a *unit cost* qRAM lookup operation.

How: bringing it all together

So in toto


- pick your lattice sieve,
- determine its operation in terms of (k, n) , d and internal sieve parameters,
- determine the quantum circuit for amplitude amplification,
- pick your cost metric for quantum computation,
- minimise the cost under chosen metric in terms of (k, n) and internal sieve parameters. . .

Estimates: ListDecoding depth-width



ListDecodingSearch. Comparing c: (RAM) with q: (depth-width), and the leading terms of the asymptotic complexities.

Estimates: ListDecoding Gidney–Ekerå error correction



ListDecodingSearch. Comparing c: (RAM) with q: (GE19), and the leading terms of the asymptotic complexities.

Our estimates suggest less advantage for this quantum sieve than the asymptotic $2^{(0.292-0.265)d+o(d)}$, without entirely ruling out their relevance.

Discussion I

Our estimates suggest less advantage for this quantum sieve than the asymptotic $2^{(0.292-0.265)d+o(d)}$, without entirely ruling out their relevance.

Quantum Metric	d	$\log_2 \text{time}_C$	$\log_2 \text{time}_Q$	asym	$\log_2 \text{memory}$
Gidney–Ekerå	312	119	119	8	78
Gidney–Ekerå	352	130	128	10	87
Gidney–Ekerå	824	270	256	22	187
Depth-Width	544	189	176	15	128
Gidney–Ekerå	544	189	182	15	128

All classical costs are in a simple RAM model, the above table is for ListDecoding.

Our analyses do not account for the cost of qRAM and RAM, required in $\mathbf{G}(g)$ and g respectively, to which we assign unit cost. Neither has unit cost in practice, but qRAM is expected to have a much higher cost.

We also do not capture the natural clock speed error correction implies: after each layer of quantum circuit depth non-trivial classical processing must occur.

Finally, we do not apply depth constraints, the impact of which on quantum search is more than classical search, which can be trivially parallelised.

The NNS search routine we cost is *not* the same as SVP.

The NNS search routine we cost is *not* the same as SVP.

Under?

- the NNS subroutine is iterated $\text{poly}(d)$ times,
- other subroutines, e.g. bucketing or lattice sampling, are not accounted for.

The NNS search routine we cost is *not* the same as SVP.

Under?

- the NNS subroutine is iterated $\text{poly}(d)$ times,
- other subroutines, e.g. bucketing or lattice sampling, are not accounted for.

Over?

- using ‘dimensions for free’ techniques [Duc18], NNS in dimension d solves SVP in dimension $d' > d$,
- many heuristic tricks [DSvW21, ADH⁺19, FBB⁺14] are not captured.

Thanks




All data and our software can be found at

<https://github.com/jschanck/eprint-2019-1161>

The paper can be found at

<https://eprint.iacr.org/2019/1161>

-  Martin R. Albrecht, Léo Ducas, Gottfried Herold, Elena Kirshanova, Eamonn W. Postlethwaite, and Marc Stevens, *The general sieve kernel and new records in lattice reduction*, EUROCRYPT, 2019.
-  Anja Becker, Léo Ducas, Nicolas Gama, and Thijs Laarhoven, *New directions in nearest neighbor searching with applications to lattice sieving*, Proceedings of the Twenty-Seventh Annual ACM-SIAM Symposium on Discrete Algorithms, 2016.
-  Anja Becker, Nicolas Gama, and Antoine Joux, *Speeding-up lattice sieving without increasing the memory, using sub-quadratic nearest neighbor search*, Cryptology ePrint Archive, Report 2015/522, 2015, <https://eprint.iacr.org/2015/522>.
-  Léo Ducas, Marc Stevens, and Wessel van Woerden, *Advanced lattice sieving on gpus, with tensor cores*, Cryptology ePrint Archive, Report 2021/141, 2021, <https://eprint.iacr.org/2021/141>.
-  Léo Ducas, *Shortest vector from lattice sieving: A few dimensions for free*, EUROCRYPT, 2018.
-  Robert Fitzpatrick, Christian Bischof, Johannes Buchmann, Özgür Dagdelen, Florian Göpfert, Artur Mariano, and Bo-Yin Yang, *Tuning gauss sieve for speed*, LATINCRYPT, 2014.

-  Craig Gidney and Martin Ekerå, *How to factor 2048 bit rsa integers in 8 hours using 20 million noisy qubits*, 2019.
-  Samuel Jaques and John M. Schanck, *Quantum cryptanalysis in the ram model: Claw-finding attacks on sike*, CRYPTO, 2019.
-  Phong Q. Nguyen and Thomas Vidick, *Sieve algorithms for the shortest vector problem are practical*, Journal of Mathematical Cryptology **2** (2008), no. 2, 181–207.